

# DataMining3

May 16, 2019

```
In [2]: # for basic operations
import numpy as np
import pandas as pd

# data visualization
import matplotlib.pyplot as plt
import seaborn as sns
plt.style.use('fivethirtyeight')

data = pd.read_csv('C:/Users/Kai/Desktop/Assignment3/data/breast-w_preprocess.csv')

In [3]: # for basic operations
import numpy as np
import pandas as pd

# data visualization
import matplotlib.pyplot as plt
import seaborn as sns
plt.style.use('fivethirtyeight')

data = pd.read_csv('C:/Users/Kai/Desktop/Assignment3/data/breast-w_preprocess.csv')

# printing the shape
print(data.shape)

# printing the head of the data
data.head()

# describing the data
data.describe()

# getting the info of the data
data.info()

# checking if the dataset contains any NULL Values
data.isnull().sum().sum()

sns.pairplot(data)
```

```
(683, 10)
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 683 entries, 0 to 682
Data columns (total 10 columns):
Clump_Thickness      683 non-null int64
Cell_Size_Uniformity 683 non-null int64
Cell_Shape_Uniformity 683 non-null int64
Marginal_Adhesion    683 non-null int64
Single_Epi_Cell_Size 683 non-null int64
Bare_Nuclei          683 non-null int64
Bland_Chromatin      683 non-null int64
Normal_Nucleoli      683 non-null int64
Mitoses              683 non-null int64
Class                 683 non-null object
dtypes: int64(9), object(1)
memory usage: 53.4+ KB
```

```
Out[3]: <seaborn.axisgrid.PairGrid at 0x1d8c5efa630>
```



In [14]: # boxen plots

```
plt.rcParams['figure.figsize'] = (18, 16)
```

```
plt.subplot(2, 2, 1)
```

```
sns.boxplot(x = data['Class'], y = data['Cell_Size_Uniformity'], data = data, palette
```

```
plt.title('Class vs Cell_Size_Uniformity', fontsize = 20)
```

```
plt.subplot(2, 2, 2)
```


```
sns.boxplot(x = 'Class', y = 'Cell_Shape_Uniformity', data = data, palette = 'summer')
```

```
plt.title('Class vs Cell_Shape_Uniformity', fontsize = 20)
```

```
plt.subplot(2, 2, 3)
sns.boxplot(x = 'Class', y = 'Marginal_Adhesion', data = data, palette = 'spring')
plt.title('Class vs Marginal_Adhesion', fontsize = 20)

plt.subplot(2, 2, 4)
sns.boxplot(x = 'Class', y = 'Single_Epi_Cell_Size', data = data, palette = 'deep')
plt.title('Class vs Single_Epi_Cell_Size', fontsize = 20)

plt.show()
```



```
In [16]: plt.rcParams['figure.figsize'] = (18, 16)
```

```
plt.subplot(2, 2, 1)
sns.violinplot(x = data['Class'], y = data['Bare_Nuclei'], data = data, palette = 'ra
plt.title('Class vs Bare_Nuclei', fontsize = 20)

plt.subplot(2, 2, 2)
```

```

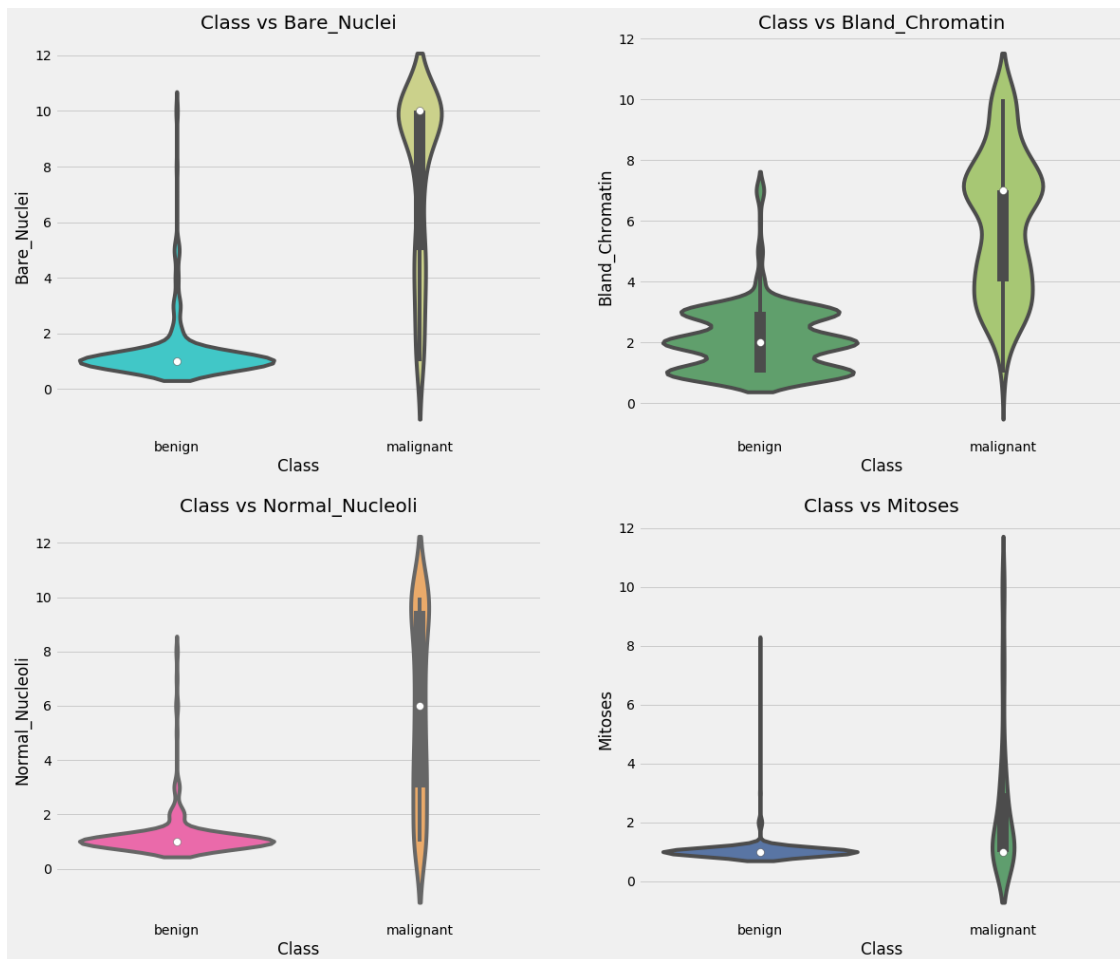
sns.violinplot(x = 'Class', y = 'Bland_Chromatin', data = data, palette = 'summer')
plt.title('Class vs Bland_Chromatin', fontsize = 20)

plt.subplot(2, 2, 3)
sns.violinplot(x = 'Class', y = 'Normal_Nucleoli', data = data, palette = 'spring')
plt.title('Class vs Normal_Nucleoli', fontsize = 20)

plt.subplot(2, 2, 4)
sns.violinplot(x = 'Class', y = 'Mitoses', data = data, palette = 'deep')
plt.title('Class vs Mitoses', fontsize = 20)

plt.show()

```



In [20]: # boxen plots

```

plt.rcParams['figure.figsize'] = (18, 16)

plt.subplot(2, 2, 1)

```

```

sns.swarmplot(x = 'Class', y = 'Marginal_Adhesion', data = data, palette = 'spring')
plt.title('Class vs Marginal_Adhesion', fontsize = 20)

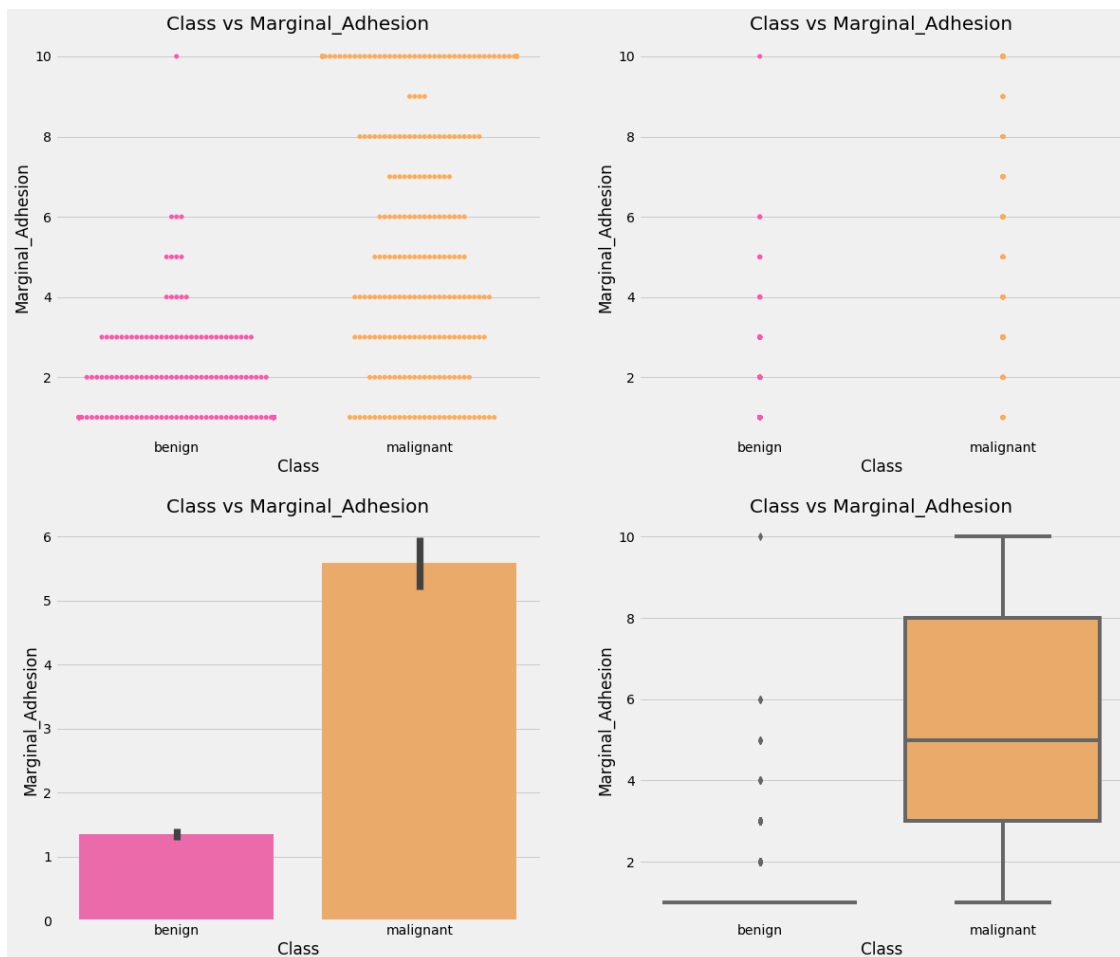
plt.subplot(2, 2, 2)
sns.stripplot(x = 'Class', y = 'Marginal_Adhesion', data = data, palette = 'spring')
plt.title('Class vs Marginal_Adhesion', fontsize = 20)

plt.subplot(2, 2, 3)
sns.barplot(x = 'Class', y = 'Marginal_Adhesion', data = data, palette = 'spring')
plt.title('Class vs Marginal_Adhesion', fontsize = 20)

plt.subplot(2, 2, 4)
sns.boxplot(x = 'Class', y = 'Marginal_Adhesion', data = data, palette = 'spring')
plt.title('Class vs Marginal_Adhesion', fontsize = 20)

plt.show()

```



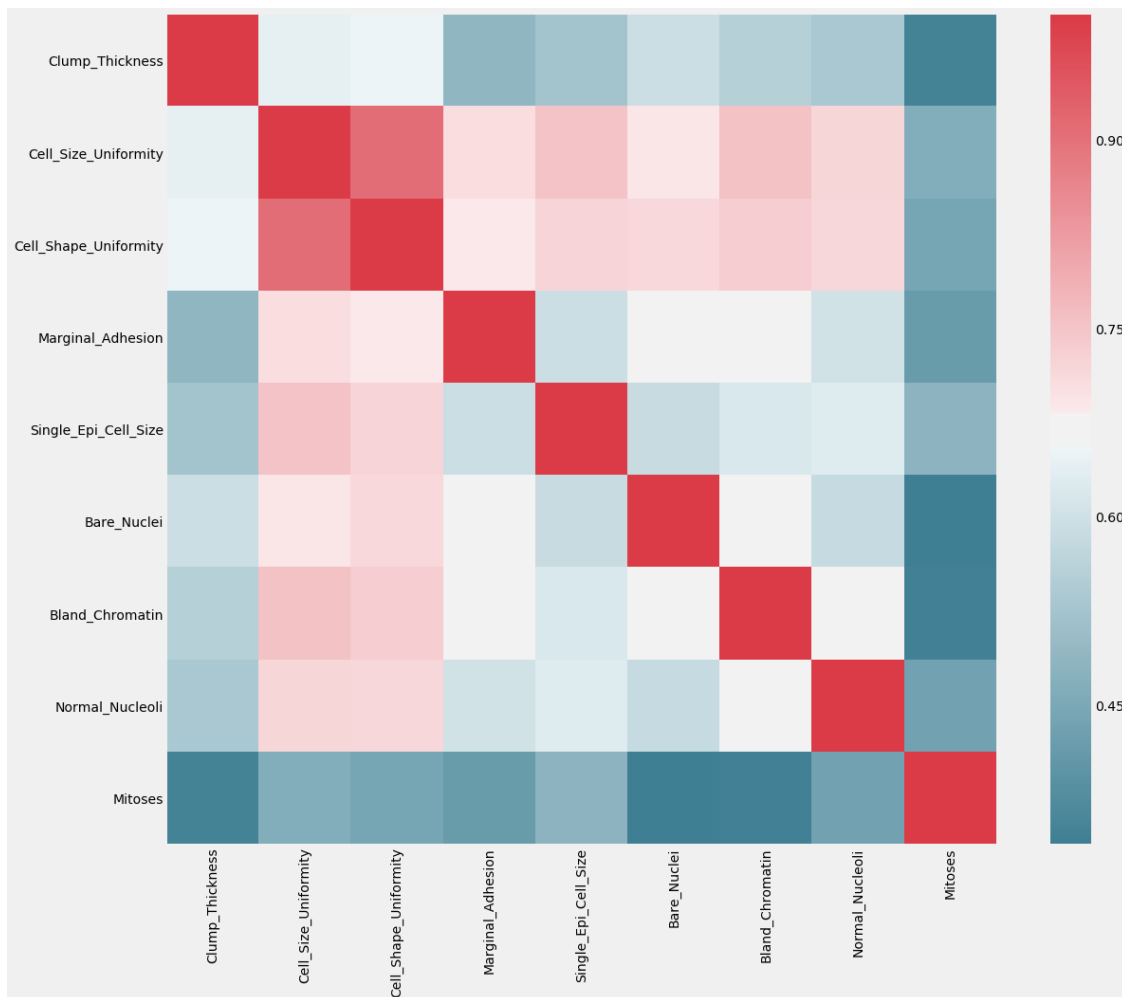
```

In [21]: f, ax = plt.subplots(figsize=(18, 15))

```

```
corr = data.corr()
sns.heatmap(corr, mask=np.zeros_like(corr, dtype=np.bool), cmap=sns.diverging_palette(
    square=True, ax=ax)
```

Out [21]: <matplotlib.axes.\_subplots.AxesSubplot at 0x1d8cbd0de10>



```
In [22]: # DATA PREPROCESSING

# label encoding of the dependent variable

# importing label encoder
from sklearn.preprocessing import LabelEncoder

# performing label encoding
le = LabelEncoder()
data['Class'] = le.fit_transform(data['Class'])
```

```
In [23]: data['Class'].value_counts()
```

```
Out[23]: 0    444
         1    239
         Name: Class, dtype: int64
```

```
In [65]: # splitting the dependent and independent variables from the dataset
```

```
x = data.iloc[:,0:9]
y = data.iloc[:,9]
```

```
print(x.shape)
print(y.shape)
```

```
(683, 9)
(683,)
```

```
In [70]: from sklearn.model_selection import train_test_split
```

```
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.25, random_st
```

```
print(x_train.shape)
print(y_train.shape)
print(x_test.shape)
print(y_test.shape)
```

```
(512, 9)
(512,)
(171, 9)
(171,)
```

```
In [71]: # standard scaling
```

```
from sklearn.preprocessing import StandardScaler
```

```
sc = StandardScaler()
```

```
x_train = sc.fit_transform(x_train)
x_test = sc.fit_transform(x_test)
```

```
In [90]: # LOGISTIC REGRESSION
```

```
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.model_selection import GridSearchCV, cross_val_score
```

```
# creating the model
```



```

model = DecisionTreeClassifier()

# feeding the training data into the model
model.fit(x_train, y_train)
y_prob = model.predict_proba(x_test)[: ,1] # This will give you positive class predict
# print(y_prob)
y_pred = model.predict(x_test)

# Calculating the accuracies
print("Training accuracy :", model.score(x_train, y_train))
print("Testing accuracy :", model.score(x_test, y_test))

# classification report
cr = classification_report(y_test, y_pred)
print(cr)

# confusion matrix
plt.rcParams['figure.figsize'] = (5, 5)
cm = confusion_matrix(y_test, y_pred)
sns.heatmap(cm, annot = True, cmap = 'cool')


```

Training accuracy : 1.0

Testing accuracy : 0.935672514619883

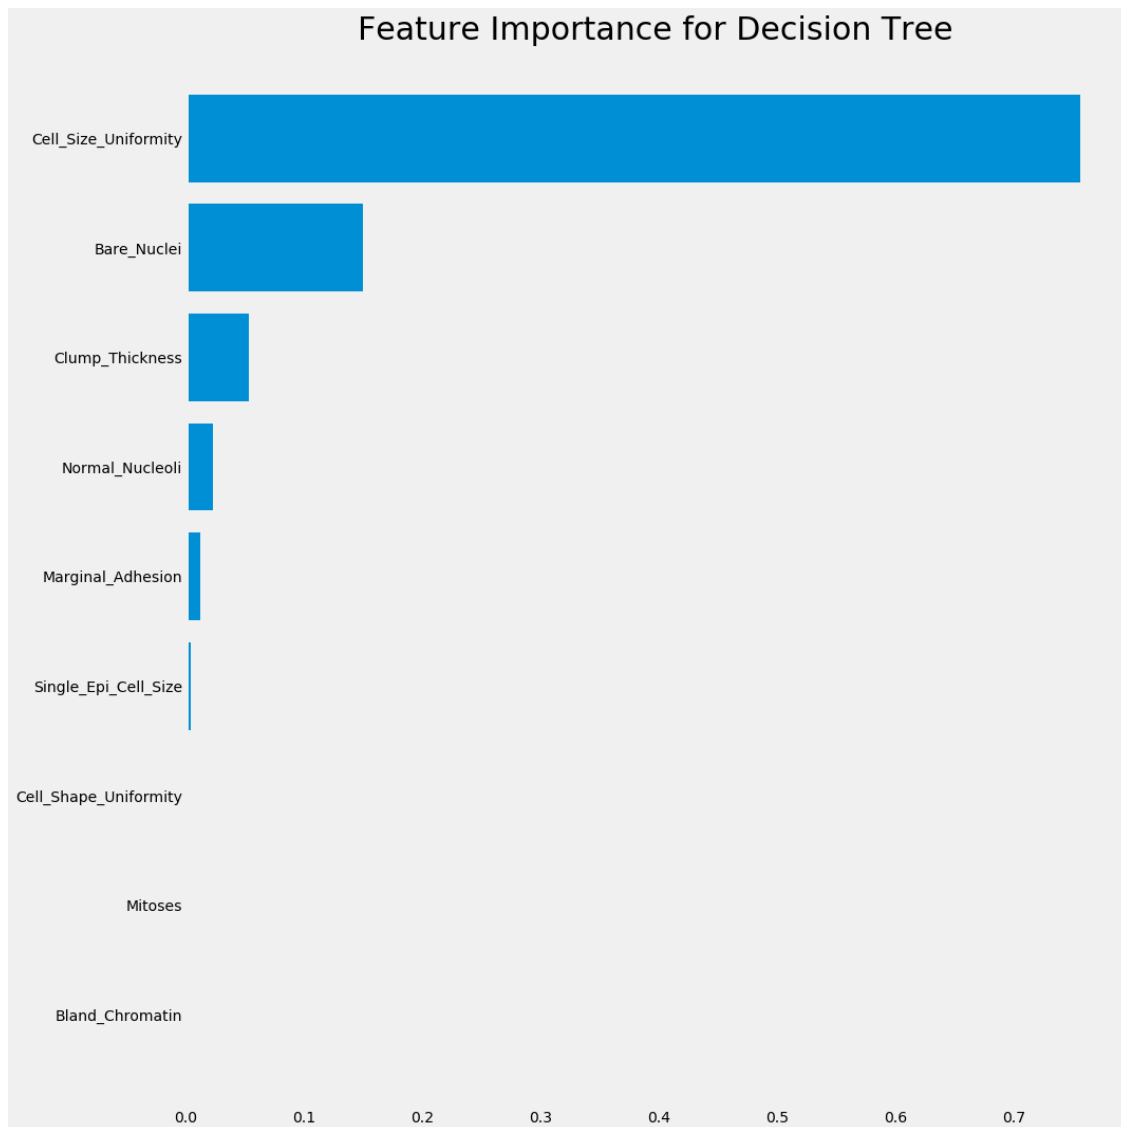
	precision	recall	f1-score	support
0	0.95	0.96	0.95	117
1	0.91	0.89	0.90	54
avg / total	0.94	0.94	0.94	171

Out[90]: <matplotlib.axes.\_subplots.AxesSubplot at 0x1d8cc0a22b0>



```
In [91]: features = data.columns
         importance = model.feature_importances_
         indices = np.argsort(importance)

         plt.rcParams['figure.figsize'] = (15, 15)
         plt.barh(range(len(indices)), importance[indices])
         plt.yticks(range(len(indices)), features[indices])
         plt.title('Feature Importance for Decision Tree', fontsize = 30)
         plt.grid()
         plt.tight_layout()
         plt.show()
```



```
In [92]: from sklearn import metrics
auc_roc=metrics.roc_auc_score(y_test,y_pred)
print("auc_roc: ")
print(auc_roc)

from sklearn.metrics import roc_curve, auc
false_positive_rate, true_positive_rate, thresholds = roc_curve(y_test, y_prob)
roc_auc = auc(false_positive_rate, true_positive_rate)
print("roc_auc: ")
print(roc_auc)

import matplotlib.pyplot as plt
plt.figure(figsize=(10,10))
```

```

plt.title('ROC curve_Decision Tree')
plt.plot(false_positive_rate,true_positive_rate, color='red',label = 'AUC = %0.2f' % auc)
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1],linestyle='--')
plt.axis('tight')
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')

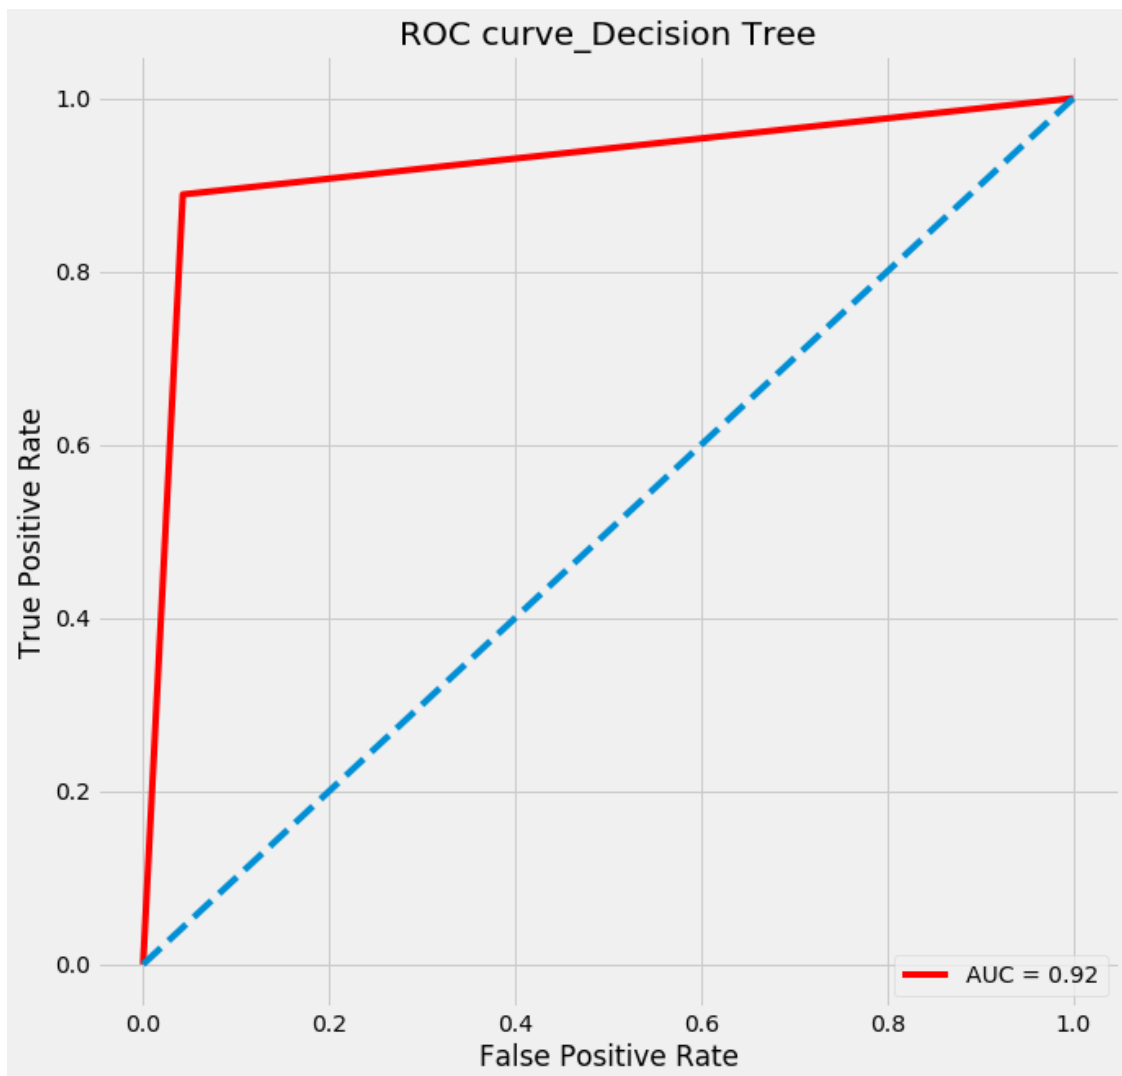
```

```

auc_roc:
0.9230769230769231
roc_auc:
0.9230769230769231

```

Out[92]: Text(0.5,0,'False Positive Rate')



In [96]: *# Support Vector Machine*

```
from sklearn.svm import SVC

# creating a model
model = SVC()

# feeding the training data into the model
model.fit(x_train, y_train)

# y_prob = model.predict_proba(x_test)[: ,1] # This will give you positive class prediction

# predicting the test set results
y_pred = model.predict(x_test)

# Calculating the accuracies
print("Training accuracy :", model.score(x_train, y_train))
print("Testing accuracy :", model.score(x_test, y_test))

# classification report
cr = classification_report(y_test, y_pred)
print(cr)


# confusion matrix
plt.rcParams['figure.figsize'] = (5, 5)
cm = confusion_matrix(y_test, y_pred)
sns.heatmap(cm, annot = True, cmap = 'rainbow')
```

Training accuracy : 0.982421875

Testing accuracy : 0.9590643274853801

	precision	recall	f1-score	support
0	0.98	0.96	0.97	117
1	0.91	0.96	0.94	54
avg / total	0.96	0.96	0.96	171

Out[96]: <matplotlib.axes.\_subplots.AxesSubplot at 0x1d8cd974128>



```
In [94]: from sklearn import metrics
auc_roc=metrics.roc_auc_score(y_test,y_pred)
print("auc_roc: ")
print(auc_roc)

from sklearn.metrics import roc_curve, auc
false_positive_rate, true_positive_rate, thresholds = roc_curve(y_test, y_prob)
roc_auc = auc(false_positive_rate, true_positive_rate)
print("roc_auc: ")
print(roc_auc)

import matplotlib.pyplot as plt
plt.figure(figsize=(10,10))
plt.title('ROC curve_Decision Tree')
plt.plot(false_positive_rate,true_positive_rate, color='red',label = 'AUC = %0.2f' % roc_auc)
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1],linestyle='--')
plt.axis('tight')
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
```

```
auc_roc:  
0.9601139601139601  
roc_auc:  
0.9230769230769231
```

```
Out[94]: Text(0.5,0,'False Positive Rate')
```

